2. Exercises: Concurrency in Java and its memory model

DCC-FCUP, University of Porto
José Proença
Concurrent Programming – Part 2

These exercises are taken mainly from the book "*Learning Concurrent Programming in Scala*", and are designed to test the knowledge of the Scala programming language. You are required to implement higher-level concurrency abstractions in terms of basic JVM concurrency primitives. The exercises are not ordered in any particular order, but some assume earlier exercises have been done.

Exercise 1. Implement a parallel method, which takes two computation blocks, a and b, and starts each of them in a new thread. The method must return a tuple with the result values of both the computations. It should have the following signature:

```scala
def parallel[A, B](a: =>A, b: =>B): (A, B)
```

Place the thread and the log function, defined in the theoretical lessons, in a package object and use them when defining parallel.

Exercise 2. Implement a periodically method, which takes a time interval duration specified in milliseconds, and a computation block b. The method starts a thread that executes the computation block b every duration milliseconds. It should have the following signature:

```scala
def periodically(duration: Long)(b: =>Unit): Unit
```

Extra: try to use a daemon thread.

Exercise 3. Implement a SyncVar class with the following interface:

```scala
class SyncVar[T] {
  def get(): T = ???
  def put(x: T): Unit = ???
}
```

A SyncVar object is used to exchange values between two or more threads. When created, the SyncVar object is **empty**:

- calling get throws an exception, and

- calling put adds a value to the SyncVar object.

After a value is added to a SyncVar object, we say that it is **non-empty**:

- calling get returns the current value, and changes the state to empty, and

- calling put throws an exception.

Exercise 4. The SyncVar object from the previous exercise can be cumbersome to use, due to exceptions when the SyncVar object is in an invalid state. Implement a pair of methods, isEmpty and nonEmpty, on

the SyncVar object. Then, using these, implement a **producer thread** that transfers a range of numbers 0 until 15, and a **consumer thread** that receives these (using busy-waiting) and prints them. Run both threads in parallel.

Exercise 5. Using the isEmpty and nonEmpty pair of methods from the previous exercise requires busy-waiting. Add the following methods to the SyncVar class:

```
def getWait(): T
def putWait(x: T): Unit
```

These methods have similar semantics as before, but go into the waiting state instead of throwing an exception, and return once the SyncVar object is empty or non-empty, respectively.

Exercise 6. A SyncVar object can hold at most one value at a time. Implement a SyncQueue class, which has the same interface as the SyncVar class, but can hold at most n values. The n parameter is specified in the constructor of the SyncQueue class.

Exercise 7. The send method in the Deadlocks' slides was used to transfer money between the two accounts. The sendAll method takes a set accounts of bank accounts and a target bank account, and transfers all the money from every account in accounts to the target bank account. The sendAll method has the following signature:

```
def sendAll(accounts: Set[Account], target: Account): Unit
```

Implement the sendAll method and ensure that a deadlock cannot occur.

Exercise 8. Recall the asynchronous method from the Guarded blocks' slides. This method stores the tasks in a First In First Out (FIFO) queue; before a submitted task is executed, all the previously submitted tasks need to be executed. In some cases, we want to assign priorities to tasks so that a high-priority task can execute as soon as it is submitted to the task pool. Implement a PriorityTaskPool class that has the asynchronous method with the following signature:

```
def asynchronous(priority: Int)(task: =>Unit): Unit
```

A single worker thread picks tasks submitted to the pool and executes them. Whenever the worker thread picks a new task from the pool for execution, that task must have the highest priority in the pool.

[**Suggestion:** *use the scala.collection.mutable.PriorityQueue data structure.*   ]

Exercise 9. Extend the PriorityTaskPool class from the previous exercise so that it supports any number of worker threads p. The parameter p is specified in the constructor of the PriorityTaskPool class.

Exercise 10. Extend the PriorityTaskPool class from the previous exercise so that it supports the shutdown method:

```
def shutdown(): Unit
```

When the `shutdown` method is called, all the tasks with the priorities greater than `important` must be completed, and the rest of the tasks must be discarded. The `important` integer parameter is specified in the constructor of the `PriorityTaskPool` class.

Exercise 11. Implement a `ConcurrentBiMap` collection, which is a concurrent bidirectional map. The invariant is that every key is mapped to exactly one value, and vice versa. Operations must be atomic. The concurrent bidirectional map has the following interface:

```
class ConcurrentBiMap[K, V] {
  def put(k: K, v: V): Option[(K, V)
  def removeKey(k: K): Option[V]
  def removeValue(v: V): Option[K]
  def getValue(k: K): Option[V]
  def getKey(v: V): Option[K]
  def size: Int
  def iterator: Iterator[(K, V)]
}
```

Make sure that your implementation prevents deadlocks from occurring in the map.

Exercise 12. Add a `replace` method to the concurrent bidirectional map from the previous exercise. The method should atomically replace a key-value pair with another key-value pair:

```
def replace(k1: K, v1: V, k2: K, v2: V): Unit
```

Exercise 13. Test the implementation of the concurrent bidirectional map from the earlier exercise by creating a test in which several threads concurrently insert millions of key-value pairs into the map. When all of them complete, another batch of threads must concurrently invert the entries in the map – for any key-value pair `(k1, k2)`, the thread should replace it with a key-value pair `(k2, k1)`.